

Diamond Software Documentation

Version SNAPSHOT

Table of Contents

1. Introduction.....	2
1.1. Purpose	2
1.2. Decision Documentation.....	2
2. Architecture.....	3
2.1. Service Types	3
2.1.1. Monolith	3
2.1.2. Modulith	3
2.1.3. Macroservices	3
2.1.4. Microservices	3
2.1.5. Nanoservices	4
2.2. Layering	4
2.2.1. Subprojects.....	5
2.2.2. Client-SDKs	6
2.3. Technology Choices	7
2.3.1. ADR - Language Choice	7
2.3.2. ADR - Webframework Choice.....	7
2.3.3. ADR - Persistence Technology	8
3. API Design	9
3.1. ADR - Filter ReST Resources	9
3.2. ADR - Generate Custom OpenAPI Code	10
3.2.1. Resources	10
3.3. ADR - Paginate ReST Resources.....	10
3.4. ADR - Secure HTTP Endpoints.....	12
3.5. ADR - Sort ReST Resources	12
3.6. ADR - Versioning of HTTP Endpoints.....	13
4. Application.....	15
4.1. Tech Stack.....	15
4.2. ADR - Application Configuration.....	15
4.3. ADR - Layer Granularity	16
4.4. KPIs	16
4.5. Startup Arguments.....	16
4.6. Gradle	17
4.6.1. Gradle Properties	17
4.6.2. Gradle Tasks.....	17
4.6.3. ADR - Manage Dependencies	18
4.6.4. ADR - Use Gradle Kotlin DSL.....	20
4.7. Documentation	20
4.7.1. ADR - Documentation Tool	20

4.7.2. ADR - Diagram Drawing Tool	21
5. Coding	22
5.1. Coding Philosophy	22
5.1.1. Best Practices	22
5.2. ADR - Generating Sources	23
5.3. Local Development	24
5.3.1. IntelliJ Plugins	24
5.4. Libraries	25
5.4.1. ADR - Database Access Library	25
5.4.2. ADR - Scheduler Library	26
5.4.3. ADR - Access SFTP service	27
5.4.4. ADR - DateTime library	27
5.4.5. ADR - Logging with Kotlin	28
5.4.6. ADR - Logging Configuration	29
6. Testing	30
6.1. Test Types	30
6.1.1. Unit Tests	30
6.1.2. Integration Tests	30
6.1.3. End-to-End (E2E) Tests	30
6.1.4. Other Types	31
6.2. Test Strategy	31
6.3. E1E Test	33
6.4. Best Practices	33
6.5. Anti Patterns	33
6.5.1. Testable Code	33
6.5.2. Maintainable Tests	33
6.6. Test Tech	34
6.6.1. ADR - Use BDD	34
6.6.2. ADR - Use Cucumber Lambdas	34
6.6.3. ADR - Verify Coverage	34
6.6.4. ADR - Assert JSON	35
7. Code Quality	36
7.1. Static Code Analysis	36
7.2. SonarQube	36
7.3. ADR - Static Code Analysis	36
7.4. ADR - Enforce Failure Handling	36
8. DevOps	38
8.1. GitHub Actions	38
8.2. Environment Variables	38
8.3. Distributed Services	39
8.3.1. ADR - Versioning Scheme	39

9. Appendix	40
9.1. Resources	40

If you want to go fast and far in a sustainable way, you need to go slow and lightweight with your friends.

This document was built using AsciiDoc, a simple yet powerful markup language for writing technical documentation.

The build for this document was triggered on Sat, 6. December 2025 at 21:24 - enjoy :)

Chapter 1. Introduction

1.1. Purpose

TODO

1.2. Decision Documentation

We use ADRs (Architecture Decision Records) to document important architectural and technical decisions made during the development of this software project.

For more explanation what ADRs are, please read: <https://github.com/seepick/diamond/blob/main/doc/decisions/README.adoc>

Chapter 2. Architecture

2.1. Service Types

- single codebase, or multiple (classpath isolated) codebases
- single service (pseudo-monolith) vs multiple services (isolation, performance)
- TODO read and summarize: <https://nordicapis.com/whats-the-difference-microservice-macro-service-monolith/>

2.1.1. Monolith

- one codebase, one deployable
- GOOD: low complexity (versioning, service location) and high speed (calls via code in same process instead network)
- BAD: change-resistant, slow, ...

2.1.2. Modulith

- one codebase, basically just a monolith but a bit cleaner packaging; domain wise, not tech
- all the disadvantages of a monolith
- degrees:
- low version: only java packages (enforce package import; use spring modulith)
- mid version: maven sub-modules (sub-projects in gradle) or external repos with dependency on them
- high version: classpath isolated libraries (OSGi bundles, plugin architecture; extreme low coupling)
- this seems to be the only one counteracting the monolith's nature of low/mid version

2.1.3. Macroservices

- bigger sized microservices
- a bit less (manageable) complexity (pros of monolith) and still good enough to gain from the benefits (pros of modulith)

2.1.4. Microservices

- one very small codebase to one very small deployable
- GOOD:
- one service down, rest of application still runs
- very quick build time; code very isolated, small, thus easy to change
- BAD:

- complexity, service orchestration, versioning

2.1.5. Nanoservices

- deploy functions (AWS lambda)
- seems too extreme; complete different approach
- <https://medium.com/@shrinit.poojary12/microservices-vs-nano-services-a-detailed-comparison-993453e65b8c>

2.2. Layering

Usecase: A feature **change** in subdomain **X**




Figure 1. The degree of impact on the application depending on the choice of packaging strategy

Given the usecase of having to add a new field to a specific domain. All three layers need to be touched: view, logic, and persistence for that domain.

Three approaches are possible:

1. Technical Packaging

- This is what we usually do; down to the extreme of packages like exception or even enum.
- In terms of single responsibility this is doubtful (simply looking at the import statements).
- When touching one aspect, all layers (with all domains), basically the whole application was changed (no caching/reusing of already built modules).

2. Domain Packaging

- Java's package/namespace idea of reversed DNS names was initially intended for domain packaging.
- E.g. When a feature login needs to be implemented, only classes within the domain package auth will be touched (view, logic, and persistence in one package).
- But now the whole classpath is polluted with all kind of different technical aspects (view

controller next to DB repositories)...

3. Combined Approach

- To split it up and we gained the best of both worlds; with some added complexity.

Final thought:

- This principle not only applies to packages, but also to services; in the big and in the small.
- There is no absolute right way of doing it; any solution will have pros and cons.
- Often the vertical slice is represented as a microservice, thus within it the domain is already split (it could be split even further though; too much?), thus only tech packaging relevant.
 - Tech split within domain split.
- And always then the need (if no horizontal layer) to create a shared library for each layer...

2.2.1. Subprojects

Each architectural layer is represented by a Gradle sub-project (or sub-module as called in Maven terminology).

Main Subprojects:

- app
- view
 - view-routing and view-controller-[api/impl] could be merged into one...
- logic-api
- logic-impl
- extern
 - extern-api
 - extern-impl (HTTP, SFTP)
 - extern-stub
 - extern-model (?) - generated source code (YAML, WSDL)
- persistence
 - persistence-api
 - persistence-impl (is actually persistence-exposed)
 - once there was a persistence-stub, but it turned out to be useless (already in full control)

Other Subprojects:

- client-sdk (kotlin MP)
 - sdk
 - models
 - tests?

- openapi-gen (resides in shared for now, needs to be externalized)

Test Subprojects:

- itest
- etest: separate GIT repo (maybe dep on client-model)

Shared Subprojects:

- test
- logging

2.2.2. Client-SDKs

What is it?

- A (provided) client library to integrate a service (backend, "extern")
- It contains:
 - DTO models; serializable data transfer object.
 - API Client interfaces; either ready-to-use controllers or simply header interfaces.
 - Reusable wiremock preparations (contract alignment)
 - Reusable test instances for property based testing.

Why?

- It reduces integration-effort (code to write and maintain).
- It takes risk of writing integration tests based on wrong assumptions (contract).
 - Capturing behavior in a formal way which can be tested against.

Code generation:

- Preferably most is generated by an OpenAPI generator (contract enforcement).
 - Not everything is doable with code generation (at least not with realistic effort); some needs to be handwritten.
- It is preferably available in all used client languages.
 - Different client technologies exist; in different versions.
- Alternative: Provide custom generators to be used.

Best Practices:

- Every service should provide such a SDK; if not, implement ourselves.
 - The code resides in its own repository; separately built and versioned.
- We maintain a client-useful (confirmed benefit) SDK for our API ourselves.
 - We potentially used it ourselves in tests, http-controllers, etc.

2.3. Technology Choices

JVM based, not .NET based, because "Java is dead, long live Java". The (Java) ecosystem is fantastic, opensource, mature, and stable. The (Java) programming language is old and outdated. Other languages are nowadays running on the JVM, like Kotlin and Scala. Especially Kotlin is fully interoperable with Java and thus all the tools can be reused and for developers a smooth migration path exists.

2.3.1. ADR - Language Choice

- Status: Draft
- Created: 5.12.2025
- Author: Seepick

Java or Kotlin?

TODO

Definitely no Scala, Clojure, ... and no non-JVM languages (Go, Rust, Python, C#, ...).

EDU: from plain java, to Lombok, to kotlin (data classes, properties, prop-initializing-ctor; from final, to var, to val)

EDU: from plain java, to instancio, to kotest-property

2.3.2. ADR - Webframework Choice

- Status: Draft
- Created: 5.12.2025
- Author: Seepick

Spring Boot or Ktor? (Quarkus?)

Context:

- besides people dont know it, less mature, less doc
- much needs handwritten (openapi generation, jpa orm)
- the thing with spring boot is, very much like maven, "super strict, either our or no way"; very opinionated (convention over configuration like)
 - it's beneficial, speed up, and stuff, but once you want to do more custom made things, it's a fucking pain in the ass (like openshift: you are a pro user, use k8s directly)
- if you want to start only certain parts of the application in an isolated way, and then also in parallel, then spring boot is impossible for that

- trust me, i tried it several times. it's fucked.
- with ktor, as everything is code (instead of annotations), there is a beautifully simple way to do it, and so extremely flexible, that every test strategy we can imagine can be implemented.

2.3.3. ADR - Persistence Technology

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- We need to choose a persistence technology for our application.
- A clear distinction between the needs for production vs test/development (local startup).
- We choose for a relational database...
 - No NoSQL store (MongoDB) or in-memory cache

Requirements:

Production DB:

- Available in corporate environment.
- If performance (searching/filtering) is an issue, maybe an elasticsearch index could be useful?

Test/Dev DB:

- Lightweight (in-memory)
- Fast startup and execution time
- No setup costs (installing, configuration, etc)

Options:

Production:

1. Oracle
2. Postgres/MySQL
3. MS SQL Server

Test/Dev:

1. H2
2. HSQLDB
3. Sqlite

Chapter 3. API Design

- OpenAPI even needed?
 - The spec is the source of truth, the holy contract
 - Especially multiple consumers (always at least 2: 1x FE and testers)
 - OpenAPI generators can keep contract in sync with code
- Show SwaggerUI as endpoint (display both UI and yaml-spec); consumer convenience

3.1. ADR - Filter ReST Resources

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- the user (via the frontend) needs to be able to filter bigger datasetes
- different fields have different types require different operators
 - general: equal, not equal
 - numeric: greater/lesser (or equal); between is done by adding both ;)
 - enum: (with)in
- operators are all in conjunctive form (coupled by AND)

Proposal:

Simple:

- `/crystals?filter=author eq 'Peter'`
- `/crystals?author=eq:Peter`
- `/crystals?weight=gt:3&weight=lte:10`
- `/crystals?state=in:F00,BAR`

Complex:

- `/crystals?filter[0].property=owner.name&filter[0].operator>equals&filter[0].value=Peter`

Radical other approach would be to turn it into a POST (PUT?) request and provide a request body (full control). It feels like we are trying to squeeze a proper payload into query parameters (feels wrong, a smell we are misusing something).

Links:

- <https://restfulapi.net/api-pagination-sorting-filtering/>

3.2. ADR - Generate Custom OpenAPI Code

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Reason:

- We need to keep the API spec (YAML) and the code (Kotlin) in sync; we don't want to rely on doing this manually.
- Thus, establish an automatism via generation.
- We need something that fits our needs: modern, up2date, high quality code and libs.
- Current (Kotlin) generators are out of date (unusable), thus a custom generator is required.

Alternatives:

- Write some tests to verify sync of spec and code.
- No, as we want immediate feedback (compile time) about correct implementation.
- Use Java generators, they are more modern and Kotlin-Java interoperability works well.
- It would not clash as Java generated clients use different libraries (even when outdated).

Implementation requirements:

- Kotlin data classes with Kotlinx serialization annotations.
- Ability to hook specific type serializer (provided or custom).
- Think of `java.time.LocalDateTime` and enums.
- Generated separately (view-model subproject) from server routes; a unique requirement not provided by default generators.
- Interfaces for Ktor server.
- Some route providing interface types could be injected into Koin.
- Only bare skeleton mandatory (HTTP method and path), everything else optional (could go crazy with details).
- Client generation is optional.

3.2.1. Resources

- <https://www.baeldung.com/java-openapi-custom-generator>
- <https://deepwiki.com/OpenAPITools/openapi-generator/5.1-creating-custom-generators>

3.3. ADR - Paginate ReST Resources

- Status: Draft

- Created: 1.12.2025
- Author: Seepick

Approaches:

Offset:

- offset, limit; skip, take
- based upon it can implement page-based pagination (could support both actually)
- bad for big datasets; DB has to scan through rows

Others:

- Keyset pagination, seek method
- for example via `since_id` and a limit/take param
- only for IDs with auto-increment, or timestamps
- Time-based pagination
- for analytics/log monitoring only
- Cursor-based pagination
- like a bookmark, a backend-determined value not necessarily linked to any data fields (contrary to keyset)
- more complex to implement
- also return the "next cursor" in the response

Best Practices:

- clearly document pagination mechanism (openAPI, provide examples)
- don't implement pagination for: A) small datasets and B) fast changing data
- use standard names
- provide page metadata:
- total pages, current page
- total items, items in page, size (request page size)
- has more
- provide navigation `_links` (HATEOES), next/previous/first/last
- either in metadata payload
- or in header: Link: `<http://localhost:8080/api/books/paged?page=1&size=10>; rel="self"`
- reuse test logic: each paginated endpoint needs to have the same set of tests

Questions:

- either use or at least get "inspired" by: <https://github.com/perracodex/exposed-pagination>

- how strict or lenient to be?
- negative numbers (skip -1, take 0)?
- take more than existing (page > totalPages)
- restrict a maximum take/limit param?

Resources:

- TODO get inspired by: <https://github.com/perracodex/exposed-pagination>
- <https://www.merge.dev/blog/rest-api-pagination>
- <https://www.merge.dev/blog/api-pagination-best-practices>
- <https://restfulapi.net/api-pagination-sorting-filtering/>

3.4. ADR - Secure HTTP Endpoints

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

How is authentication and authorization handled?

- IAM/LDAP
- Central security auth gateway in front
 - A reverse proxy (nginx) for auth (more complex pipeline)
 - Services only receive an already verified user ID header (auth-agnostic)
 - Simplifies lots of code
 - Faster build time ("IAM policy seeding" can take a lot of time)
 - Easier test setup (happier testers)
- Separation of concerns: make the "children slimmer/lighter"
- Endpoints are configured centrally (?); the obligatory "auth matrix" ;)
- OAuth2 vs JWT?

3.5. ADR - Sort ReST Resources

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

About:

- call it sort or orderBy
- orderBy=author,title

- sort direction: asc/desc suffix
- orderBy=author asc,title desc
- or plus/minus: orderBy=+author,-title
- default to asc if not defined

Error Handling:

E.g.:

Example from a fictional API

```
{
  "error": {
    "code": "InvalidOrderByExpression",
    "message": "The property 'author' in the orderby expression is not sortable",
    "details": [
      {
        "code": "UnsupportedSortProperty",
        "target": "author",
        "message": "Sorting by 'author' is not supported. Supported sort properties
are: ['id', 'title']"
      }
    ],
    "target": "orderby",
    "internal": {
      "trace-id": "00-d24f899d9c8a5a428fddd39399e7f58e-5c8a8949fcdd9a42-01",
      "timestamp": "2024-02-20T15:30:45.123Z",
      "request-id": "123e4567-e89b-12d3-a456-426614174000"
    }
  }
}
```

Links:

- <https://restfulapi.net/api-pagination-sorting-filtering/>

3.6. ADR - Versioning of HTTP Endpoints

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Options:

1. using header: Accept: application/vnd.diamond.crystal.v1+json and Content-Type: ...
 - Keeps the URLs clean; metadata in headers
 - But not so obvious at first, requires additional logic to be implemented

2. encode in path: GET /api/v1/crystals
 - out-of-the-box supported by caching

Also:

- deprecate (TODO: tag OpenAPI spec as such, generate warning in code possible?!) and communicate
 - Keep track (logging) of usage to know when to turn off

Links:

- <https://dev.to/abhivyaktii/api-versioning-using-headers-a-clean-and-flexible-approach-218p>

Chapter 4. Application



Figure 2. A precious material

4.1. Tech Stack

- Gradle
- Kotlin
- Ktor
- Koin (not really DI, but a service locator, but ok ;))
- Kotest
- Exposed
- Liquibase
- Hoplite - typesafe loading application configuration (sys-props, env-vars, files)
- Arrow
 - Especially Either for better exception handling

Clean Dependencies:

- Non-prod dependencies (test, stubs, localdev, etc like H2) should NOT be in production code
 - Instead: Activate profiles which customize the build (dependency tree)
 - E.g. for tests (presistence-tests and itests) different wiring

4.2. ADR - Application Configuration

- Status: Draft
- Created: 1.12.2025
- Author: Seepick
- config as ENV vars or centralized config manager?

Links:

- <https://softwarepatternslexicon.com/kotlin/microservices-design-patterns/configuration-management/>

4.3. ADR - Layer Granularity

- Status: Draft
- Created: 5.12.2025
- Author: Seepick

Should the view-routing and view-controller-* modules be separate or merged?

- if testing strategy is testing from http (ktor test engine) then it can be merged.
 - if testing from controller (no ktor involved, plain code, simpler but less realistic/higher risk, require complementary tests and integration of those layers) then it should be separate.
- current situation feels like adding a costly feature without using its benefit
 - the modules are separated and the itests are not using the controllers (but HTTP ktor test engine)

Separate:

- very clean
 - view-route only purpose interact with web framework and leave
 - easy to test routes
- adds complexity (already heavy due to route - controller - domain - persist/extern)
 - controller interface, controller impl, register koin, ...
 - but enforces clean architecture!
- itests could use view-controllers instead using HTTP (from cucumber looks the same, just faster/easier setup)

Merged:

- faster development; less code/complexity
- unit (integration, as they use ktor) test for routes will be heavier.

4.4. KPIs

- build time below 3mins locally (build server?)
- response times (load tests)
- lead time

4.5. Startup Arguments

- `java -jar diamond.jar printConfigOnly` - will print the read configuration object for verification
 - Used in `./local/test_apponfig.sh` to verify output

```
EnvConfig(
  ktor=KtorConfig(port=12),
  database=DatabaseConfig(
    jdbcUrl=db_url,
    username=db_user,
    password=****
  ),
  extern=ExternConfig(
    postsServiceBaseUrl=postsUrl,
    sftp=SftpConfig(
      remoteHost=sftpHost,
      port=22,
      username=sftpUser,
      authIsPassword=true,
      authPasswordOrPrivateKeyPath=****,
      knownHostsFilePath=knownHosts,
      strictHostChecking=true
    )
  )
)
```

4.6. Gradle

- Guiding Principles:
 - Keep build log output silent; don't pollute it; don't be verbose (by default)
 - Keep all warnings down to zero immediately (0-warning policy applies here as well).
- Sharing build logic via the `/buildSrc/` directory.
 - See: https://docs.gradle.org/current/userguide/sharing_build_logic_between_subprojects.html

4.6.1. Gradle Properties

with certain flags (enabled via gradle properties and on CI), it will enable more tests/reports (which take longer and require more local setup)

Pass them either as system variables `-D` (as gradle properties sometimes doesn't work `-P`).

- `diamond_version` - application version, defaults to `0` if not specified
 - displayed in info endpoint, used in generated documentation
- `runTestcontainersTests` - see task below
- `runEtests` - see task below

4.6.2. Gradle Tasks

- Build an executable JAR:

- `./gradlew :app:buildFatJar -Pdiamond_version="42"`
- Check for outdated dependencies:
- `./gradlew dependencyUpdates`
- Deploy container image to local Docker registry.
- `./gradlew publishImageToLocalRegistry`

Test:

- Run tests using testcontainers (requires Docker daemon to be running)
 - `./gradlew test -PrunTestcontainersTests=true`
- Run Karate end-to-end tests.
 - `./gradlew test -PrunEtests=true`

Documentation:

- Generate documentation for available ENV properties:
 - `./gradlew :app:generateConfigDoc`
- Generate software documentation as PDF:
 - `./gradlew :doc:SoftwareDocument:asciidoctorPdf`
- Generate software documentation as HTML:
- `./gradlew :doc:SoftwareDocument:asciidoctor`

4.6.3. ADR - Manage Dependencies

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Which way to properly manage gradle dependencies?

Problem description:

- versions repeatedly being declared, danger of mismatch (same library, different versions)
 - especially true for multi-module projects
 - leading to runtime error due to binary incompatibility (`NoSuchMethodError`, etc.)
- cumbersome as in lots to type and repeat (we want to be concise and fast)

Requirements:

- not abstracting/hiding too much (e.g. one god dependency), but be explicit/transparent enough for understanding
 - be concise enough, while capturing the essence for clarity (intuitive comprehension) sake
- not too complex, no over-engineering, just get it "good enough"

- deal with a huge amount of versions/dependencies/plugins
- custom plugin declaration (to share common code as extensions/compositions) can use it as well

Alternatives:

Either: 1) buildSrc custom code or 2) use gradle's standard version catalog.

Custom buildSrc Code:

Pros:

- it's plain kotlin code with all its advantages
 - refactoring safe
 - immediate feedback from compiler if something is wrong
 - go to source with CTRL+click (javadoc)
- full control and flexibility (custom made solution)
 - e.g. group versions and dependencies
- already used for custom plugins
 - BUT: it is NOT possible to use things from the catalog for those

Cons:

- it's custom made for something a default solution already exists
- doesn't support view usages

4.6.3.1. Gradle Version Catalog

- https://docs.gradle.org/current/userguide/version_catalogs.html

Pros:

- it's a out-of-the-box supported solution, people know it, well documented
- allows for grouped dependencies, so-called libraries

Cons:

- TOML syntax is not as useful as code (auto-completion, type feedback, etc.)
- going into source implementation/definition not possible (CTRL+click), cumbersome
- if changing/refactoring, then step-by-step errors by each build, cumbersome

Example of a Gradle version catalog

```
[bundles]
ktor = ["ktor-core", "ktor-json", "ktor-foobar"]

[plugins]
short-notation = "some.plugin.id:1.4"
```

```
versions = { id = "com.github.ben-manes.versions", version.ref = "manes-versions" }
```

Conclusion:

- both support grouping and a nested-hierarchy to control a vast amount of declarations
- both support auto-completion
- custom plugin declaration is fucked
 - neither the custom code approach, nor the version catalog supports it (too early in build phase i assume)

4.6.4. ADR - Use Gradle Kotlin DSL

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context: * obviously not groovy (cumbersome) * as project is already in kotlin, makes sense to configure gradle also with kotlin * has become stable enough the last years (support, plugins/configuration, sufficient documentation on the internet)

4.7. Documentation

4.7.1. ADR - Documentation Tool

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Options:

- Markdown
 - Well known and well supported
 - Might not provided sufficient functionality for a more complex/bigger documentation project
- Asciidoc
 - Somewhere between Markdown and LaTeX in usability/support and functionality/complexity
 - GitHub can also render it, just like markdown, yay
- LaTeX
 - Great for writing books, but definitely poses a too high barrier of entry for people unfamiliar with it
 - Local latex installation is far more than needed compared to provided infrastructure for

4.7.2. ADR - Diagram Drawing Tool

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Requirements:

- programmatic (source code, version trackable)
- but high barrier entry
- no WYSIWYG, thus not suitable for complex (and "pretty") diagrams
- defined along adoc
- easy access (no visio)
- user/rights/auth
- web based

Options:

- plantuml
- mermaid
- ...?

PlantUml:

Install dot binary from <https://graphviz.org/download/> to be able to build diagrams with code.

```
[plantuml,inlineUml,svg]
....
@startuml
[Inline] --> Interface1
[Inline] -> Interface2
@enduml
....

Or included:

[plantuml,includedUml,svg]
....
include::includes/diagram.puml[]
....
```

- With markdown (GitHub) also PlantUML possible "somehow", but not really...
 - <https://gist.github.com/noamtamim/f11982b28602bd7e604c233fbe9d910f>

Chapter 5. Coding

Kotlin.

- list the bad things as well; honesty, exception mgmt, credibility.
- intellij support; crash; young

5.1. Coding Philosophy

- fail early, fail fast
- e.g. if application config something is configured wrong/missing
- validate data (front incoming HTTP, back external systems/DB, immediately throw)
- code first / code over config
 - no strings/annotations, no declarative xml/json/yaml/properties
 - just code, be in control, flexible, safe (compiler reference check)
- no classpath scanning, reflection, other look-ups/service-locator, no auto-magically something
- avoid intrusive frameworks (the systems serves us, we don't serve the system)
- functional programming
- pure functions: stateless, side-effect free
- code as named expressions
- prefer single-expression-method `fun foo() = doSome().doOther().also { done(it) }`
- absolutely no local files (IDE specific, build generated, ...)
 - not checked in, even existing put on .gitignore
 - all auto-configured by build system/gradle
 - don't interfere with each other; give devs freedom - no coupling
- single responsibility (of package, class, method/function)
- clear names, self-explanatory, capturing its essence (single responsibility)
 - no abbreviations, no short cryptic names

5.1.1. Best Practices

- Strive for functional, pure functions (no side effects).
- instead of regular methods → "named expressions"
 - functions as named codeblocks, lead to shorter (easier to comprehend) methods
 - always use "=" after method signature; single expression body
 - no state, no "val" keyword, all just immediate execution (more functional)
 - forces to write very short methods, easy to read and comprehend
 - methods have names, they create logical structure of arbitrary code

```
fun `BAD` single expression`(): String = "foo"

fun `GOOD` single expression`() = "foo"

fun `BAD` single factory expression`(): FooDto = FooDto(
    id = 42,
)

fun `GOOD` single factory expression`() = FooDto(
    id = 42,
)

fun `BAD` db access`() {
    transactional {
        repo.update()
    }
}

fun `GOOD` db access`() = transactional {
    repo.update()
}

fun Service.`BAD` wither`(): Service {
    setFoo("bar")
    return this
}

fun Service.`GOOD` wither`() = apply {
    setFoo("bar")
}
```

- no overloaded methods, provide different name

5.2. ADR - Generating Sources

- Status: Draft
- Created: 6.12.2025
- Author: Seepick

Context:

- We use specifications to document the contract, provided by ourselves and third parties.
- Primarily OpenAPI YAML and WSDL XML files.

Should we regenerate sources each time and host the code next to our own, or should be pre-

generate and store them in a dedicated repository?

Options:

Re-generate:

- GOOD: Everything is self-contained; simple, easy.
- BAD: Slows down the build process due to regeneration despite no API changes (caching?).
- BAD: Exclusion configuration necessary for static code analysis, coverage, etc. (different rules apply to them)

Pre-generate:

- GOOD: Faster build times.
- GOOD: A change of the API contract must be a conscious act
 - Leads to more safety/stability
 - But it's also a BAD argument: slower development time, more complex!
- For 3rd party APIs, we would expect a given SDK anyways... so we do it ourselves.

Decision:

- Pre-generate!
- Divide and conquer: Manage the bigger complexity by outsourcing.
- Split things apart, separation of concerns.

5.3. Local Development

The goal is to only have to check out the sourcecode, and run a single command to start the application up in a local model. No further setup must be needed (no docker, no special files created). There is no required readme to be read, no manual or instructions; it just works out of the box.

- `run nl.uwv.smz.diamond.app.LocalDiamondApp`
- Postman collection can be found at: `/local/Diamond.postman_collection.json`
- There are a lot of useful (required) build scripts in `/local/*.sh`

5.3.1. IntelliJ Plugins

The following plugins are recommended to have a smooth development experience:

- Gradle
- EditorConfig - autoconfigures styling
- detekt - configure with `/config/detekt.yml`
- ktlint - auto-formatter

- kotest - test framework integration
- Markdown - some internal documentation files use it
- AsciiDoc - documentation written in it
- SonarQube for IDE - connector to quality gates
 - Configure with SonarQube Cloud project: https://sonarcloud.io/project/overview?id=seepick_diamond
- Gherkin - Test specifications support
- Cucumber for Kotlin - Step definitions; regexp
- Cucumber for Java - not sure, i guess it adds to it (?)
- Cucumber+ - Better support
- Cucumber Table Mapping - Specific support
- kotlin-fill-class - Useful inline code generator

Other:

- plantuml4idea (graphviz/dot required)

5.4. Libraries

5.4.1. ADR - Database Access Library

- Status: Accepted
- Created: 1.12.2025
- Author: Seepick

Context:

- We need a way to access relational database from Kotlin code.
 - We talk about DB access (as in JDBC), not a full ORM solution.
- Question is, which library to use?

Requirements:

- Needs to be type-safe (no "stringly-typed" expressions)
- Little/no code generation magic
- Kotlin-idiomatic (coroutines, DSLs)
 - For using coroutines we would need a non-blocking implementation (no JDBC)

Options:

1. Exposed by JetBrains
 - Seems production ready; mature and stable enough

- Lightweight and typesafe
2. KForm
 - ORM mapper for Kotlin, pure JDBC
 3. JPA / Hibernate
 - Heavyweight, not Kotlin-idiomatic
 - Lots of magic, code generation, proxies, etc.

Choice:

- Exposed
- It fits with the rest of the used ecosystem.
- We use the conventional DSL, not the DAO API (seem unfinished).

Example of Exposed's (unfavored) DAO API

```
// definition
class CrystalDaoEntity(id: EntityID<UUID>) : UUIDEntity(id) {
    var weightInGrams by CrystalTable.weightInGrams
    companion object : UUIDEntityClass<CrystalDaoEntity>(CrystalTable)
}
// usage
CrystalDaoEntity.findById(id)
CrystalDaoEntity.new(UUID.randomUUID()) {
    weightInGrams = 42
}
CrystalDaoEntity.findByIdAndUpdate(id) {
    it.weightInGrams = 42
}
```

5.4.2. ADR - Scheduler Library

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- we need to run jobs regularly in a cronjob way
- no spring scheduler available

Options:

- Quartz <https://www.quartz-scheduler.org>
- enterprise (heavy, slow) scheduler; the "platzhirsch", used to be the one and only; decades
- but the pros comes with cons: outdated, verbose APIs

- features: persisted in DB (requires 10+! tables), retry, clustering, prioritization etc
- no monitoring (grafana, jaeger)
- JobRunr <https://www.jobrunr.io>
- lightweight, modern; RDBMS (and NoSQL)
- features: distributed, dashboard, retry
- db-scheduler <https://github.com/kagkarlsson/db-scheduler>
- lightweight: performant, minimal deps; still offer what's need: persistence, cluster

maybe there is a mature enough kotlin library which leverages coroutines?

- KtScheduler <https://github.com/Pool-Of-Tears/KtScheduler>
- seems like a small hobby project...

Links:

- <https://www.jobrunr.io/en/blog/2024-10-31-task-schedulers-java-modern-alternatives-to-quartz/>

5.4.3. ADR - Access SFTP service

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- JSch is very old and very outdated
- Use forked and updated `com.github.mwiede:jsch` instead of original `com.jcraft:jsch`

5.4.4. ADR - DateTime library

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Use java-time or kotlin-time?

Java:

- GOOD: robust, well-known, well-integrated with libs (serialization, DB)
- BAD: cumbersome
- BAD: not supported by kotlin-idiomatic libs

Kotlin:

- GOOD: lightweight, probably/most-likely kotlin idiomatic; modern API
- GOOD: integration with kotlinx-serialization and exposed
- BAD: not all (java) libs support it
- BAD: higher barrier of entry for others
- any benefit from its multiplatform nature?

Conclusion:

- unstable API is too dangerous for production!
- NoSuchMethod error, incompatible versions on classpath
- kotlin is not well integrated yet (kotest)

5.4.5. ADR - Logging with Kotlin

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Logging evolution:

Initial approach with performance in mind

```
if (log.isDebugEnabled()) {
    log.debug("some message " + message + " was returned")
}
```

This way we only create the string if it is really necessary, bad. But it's very verbose, to do it correct and concise; bad.

Then we switched to a `String.format`-like approach to fix this, to always be able to invoke the log method without performance impact:

Less verbose, but still performant

```
log.debug("some message %s was returned", message)
```

This two issues (conciseness, performance) but still was not readable, especially when logging longer messages with multiple values. The displacement-nature of these kind of formatting approaches is just not doing it...

Finally we are doing it nicely, using Kotlin's string templates (interpolation) which are basically string concatenations, and the ease of lambda declarations (deferred/lazy evaluation):

Readable and performant

```
log.debug { "some message $dto was returned" }
```


Logger Declaration:*

Thanks to `io.github.oshaikotlin:kotlin-logging-jvm` (former "mu-kotlin") the declaration of a logger is as simple as it can be:

Implicit log factory creation

```
private val log = logger {}
```

It will look up the surrounding class and thus use the right logger factory itself. It uses Slf4j underneath.

The actual logging implementation is done by Logback (the better Log4j).

5.4.6. ADR - Logging Configuration

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- following principle: code > config
 - more flexibility, more control, more sophisticated, more robust
- programmatic logging configuration
- dev/test different from default-prod
- provide mechanism to override(?), maybe even
- TODO: runtime, to "debug" = increase log; in running prod

Chapter 6. Testing

- All tests are automated; there is not a single manual test!
- We have full confidence in our tests as a safety net.

6.1. Test Types

6.1.1. Unit Tests

- In OO-languages, the smallest unit is a class (object); tests grouped by its public methods.
- Run extremely fast, easy to write (if code is well designed/testable), stable (low maintenance).
- There is not a single framework (besides mocking if applicable) involved.
- They test business logic code, not integration code (smell: overuse of mocking).

6.1.2. Integration Tests

- Narrow definition: Everything testing 2+ classes; broad definition: Testing the application as a whole.
- Never leaves the application boundary, thus in **full control**; stable, but a bit slower.
- Relying on an (implicit/self-defined) contract is risky (better E2E).

6.1.3. End-to-End (E2E) Tests

- About:
 - Using the application as a black box; only acting upon a contract (requirements).
 - Located in a separate GIT repo (no code sharing/dependency; be independent, a parallel stream).
- Pros:
 - Verifies the (implicit) contracts
 - Tests the complexity behind configuring the systems working neatly together.
- Cons:
 - Very slow execution
 - Fragile, as depends on all systems to be up and running.
 - Sometimes the functionality is "hidden" and not testable from the outside.
 - No proper control over the (test) data.
- Details:
 - Using Karate test framework, only change base URL to target another environment.
 - Use @WIP tagging; maybe Jira ID @ABC-12345
 - Each environment (DEV, FB-*, TEST, ACC) is running them

- Could provide @ReadOnly to target also PROD
- external parties need to provide test-environment/sandbox, otherwise mock it
 - mocking external parties MUST always comply with real implementation, otherwise our (self-invented) contract is different from actual real implementation

6.1.4. Other Types

- application tests?
- module tests?
- component tests?
- acceptance/functional tests
- contract tests
- external system tests (backend)
- system tests

Other:

- ATDD, BDD Tests
- smoke tests
- regression tests
- performance/load/stress tests
 - We for sure need one of those!
 - Use gatling (with karate?) or jmeter
- security/penetration tests

Other other:

- sanity tests
- exploratory tests
- usability tests
- compatibility tests
- localization/internationalization tests

6.2. Test Strategy




Figure 3. Simplified view from different test strategies onto the application boundaries

- Some Unit Tests (using Kotest framework) covering logic (not infrastructure/IO related code).
- Lots of Integration Tests (using Cucumber framework) verifying functional requirements.
 - Either using fast test engine, or real HTTP (slower, more production-similar).
 - Either using H2 or Oracle in a testcontainer (slower, more production-similar).
 - Has the Koin context available, thus full control of the internals (extern-stub or custom mocks).
 - The parts, the IT don't cover, can be complemented by separate tests.
 - Karate not feasible anymore, due to need of custom vocabulary (step definitions).

- Few End2end Tests talking real HTTP, separately deployed, no access than a regular user would.
- A sandbox environment of 3rd party services would be great (if available).
 - Do we want to provide one ourselves to others?!

Ad complementing:

- Backend Tests, only testing the extern-layer with testcontainers.
 - Could make them system tests (without testcontainers! sandbox?), but super fragile/limited.

6.3. E1E Test

- A mix between an Integration Test and an End2end Test is a E1E Test (half of a E2E).
- It works under same assumptions as E2E (blackbox, Karate) but all real systems are mocked.
 - Additionally it has access to system internals through an additionally deployed test API.

6.4. Best Practices



Watch out to not test same multiple times, we need to be able to still move and change fast!

- Test for HTTP codes (sometimes returning 500 if authorization is not present for the user)
- Think and test blackbox (you don't know the implementation)
 - E.g. just because the same mechanism is applied to several endpoints with a cross cutting concern (security, pagination, ...), doesn't justify not testing it.
 - To overcome DRY, write sophisticated test infrastructure (Kotest include() test factory)
- There also always the possibility to deploy another service exposing additional test endpoints ;)

6.5. Anti Patterns

- Not keeping proper level of abstraction: Displaying irrelevant values.
- Hardcoded, OS-specific, absolute paths in code or properties.
- Using the filesystem; if really needed, use Java's temp file mechanism.
- Not testing "automagic" (implicit stuff; aspects, proxies; reflection, annotations)
 - Previously Spring repository interfaces (not visible in coverage as well)

6.5.1. Testable Code

- no manual instantiation or static random data generation

6.5.2. Maintainable Tests

- Assert only what's relevant, thus less in many cases.

6.6. Test Tech

6.6.1. ADR - Use BDD

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- TDD on unit test level; BDD for integration tests (on higher acceptance/functional level)
- using a gherkin-based format; use cucumber
- when having both plugins (cucumber and karate) installed, configure file types to be different
 - cucumber: *.feature
 - karate: *.karate.feature

6.6.2. ADR - Use Cucumber Lambdas

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- use more kotlin-idiomatic approach by the java8 lambda functionality
 - instead the old fashioned annotationstyle
- both approaches are fully supported by the intellij-plugin
 - click into declaration possible from feature files

Links:

- <https://github.com/cucumber/cucumber-jvm/blob/main/cucumber-java8/README.md>
- Example: <https://github.com/cucumber/cucumber-jvm/blob/main/cucumber-kotlin-java8/src/test/kotlin/io/cucumber/kotlin/LambdaStepDefinitions.kt>

6.6.3. ADR - Verify Coverage

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- Goal: automatically verify minimum test coverage threshold

- Fail build if not achieved
- Provide dashboards (Sonarqube) on current state and changes (delta)
- coverage 100%?, kotlin inlining issue...

Options:

- jacoco
 - very mature, well supported by tools
 - other coverage tools usually can export to jacoco's format too
 - not supporting Kotlin's inline
- kover
 - specific for kotlin, from jetbrains
 - bleeding edge (child sicknesses?)

6.6.4. ADR - Assert JSON

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- for itest with cucumber, we need a library which supports something like: `$.items[3].title eq 'foobar'`
- a library which we can be passed a complex string to evaluate
- for complete JSON comparison, use: <https://github.com/skyscreamer/JSONassert>

Options:

- assert full JSON: <https://github.com/skyscreamer/JSONassert>
- regular expression support: <https://fslev.github.io/json-compare/>
- jayway path is an option (`com.jayway.jsonpath:json-path`)
- it provides hamcrest matchers `com.jayway.jsonpath:json-path-assert`
- tutorial: <https://www.baeldung.com/guide-to-jayway-jsonpath>

Chapter 7. Code Quality

- quality gates
 - make the build fail upon error
 - 0-warning policy
- sonarqube (misc metrics), coverage kotlin tools (detekt, ktlint)
- CAVE: there is much overlap in the tools
 - test coverage enforcing by gradle (kover) and sonarqube
 - lots with detekt/ktlint (intellij? sonarqube?)

7.1. Static Code Analysis

Detekt

ktlint

- <https://pinterest.github.io/ktlint/1.8.0/>

7.2. SonarQube

- use the free tier for open source projects
 - See: https://sonarcloud.io/project/configuration/AutoScan?id=seepick_diamond
- define new code by days (continuous delivery), not by change delta (planned releases)
- use it to report coverage (enforced by gradle-kover itself, not sonar)

7.3. ADR - Static Code Analysis

- Status: Draft
- Created: 1.12.2025
- Author: Seepick

Context:

- actually also runtime detection possible!
- also ktlint
- enable Refactor → AutoCorrect by detekt rules
- bad: e.g. line length not auto-configured in intellij by detekt rule :-(

7.4. ADR - Enforce Failure Handling

- Status: Draft

- Created: 1.12.2025
- Author: Seepick

Context:

- Pure Functions:
 - return values are identical for identical arguments (no variation with static variables, non-local variables, mutable reference arguments or input streams, i.e., referential transparency)
 - the function has no side effects (no mutation of non-local variables, mutable reference arguments or input/output streams)
- Throwing (unchecked) exceptions is a sort of an implicit return value which can be forgotten to be handled properly, resulting in false-500 responses.
- Use the Arrow library to introduce a more functional approach, the Either class specifically
 - TODO: in the future maybe also optics to manipulate deep nested, immutable data classes.

Either:

- Good:
 - developers need to think (forced to, explicit) error types, mapping of them (translation of corrupt data to bad request/internal error)
 - supports DDD by simulating failure-aware constructors (companion operator fun with Either return type)
- Bad:
 - sometimes not clear if not familiar with it
 - sometimes IDE doesn't suggest `.bind()` although it's the solution

Chapter 8. DevOps




Figure 4. Sketching a build pipeline overview

- When a feature branch is created, a new environment gets auto-created.
 - Testers can write their E2E tests targeting it.
 - Development starts in parallel: The contract (requirements) the base.
- On merge: Old feature-branches are deleted and custom environment destroyed.

8.1. GitHub Actions

- Continuous
- Nightly
- Release

8.2. Environment Variables

Param	Type	Default	Description
DATABASE_JDBCURL	string	-	JDBC driver URL
DATABASE_PASSWORD	string	-	DB password
DATABASE_USERNAME	string	-	DB username
EXTERN_POSTSERVICEBASEURL	string	-	Base URL for the external posts service
EXTERN_SFTP_AUTHISPASSWORD	boolean	-	Whether using password or private key.
EXTERN_SFTP_AUTHPASSWORDORPRIVATEKEYPATH	string	-	Either password or private key.
EXTERN_SFTP_KNOWNHOSTSFILEPATH	string	-	Path to SSH known hosts file.
EXTERN_SFTP_PORT	integer	22	Port of the SFTP server.
EXTERN_SFTP_REMOTEHOST	string	-	Host of the SFTP server.
EXTERN_SFTP_STRICTHOSTCHECKING	boolean	true	Disable security check; do NOT activate in production; pretty please.

Param	Type	Default	Description
EXTERN_SFTP_USERNAME	string	-	Login username.
KTOR_PORT	integer	8080	Webserver HTTP port to listen to

8.3. Distributed Services

Their specifics require:

- Versioning, backwards compatiability.
- Resilience and Fault Tolerance: circuit breakers, retries, and bulkheads.

8.3.1. ADR - Versioning Scheme

- Status: Draft
- Created: 6.12.2025
- Author: Seepick

Context:

- Deployed software has to be uniquely identifiable.
- The version number needs to be generated automatically (CI/CD).
- It will be exposed in various places (APIs, UIs, documentation, etc.).

Options:

1. One single continuous version number for each build.
 - The number becomes meaningless, but very simple to implement.
2. Semantic Versioning (SemVer)
 - Well known and widely adopted.
 - Major, Minor, Patch conveys information about backward compatibility.
3. Calendar Versioning (CalVer)
 - Easy to understand and communicate.
 - Conveys information about the release date.

Resources:

- E.g. use GitHub action for that: <https://github.com/marketplace/actions/version-increment>

Chapter 9. Appendix

9.1. Resources

- ADR: <https://github.com/joelparkerhenderson/architecture-decision-record>
- SFTP inspirations:
 - <https://github.com/alkaphreak/marstech-sftp/blob/main/src/main/kotlin/fr/marstech/mtsftp/service/SftpFileConnectorServiceImpl.kt>
 - <https://medium.com/whozapp/sftp-test-implement-of-jsch-with-kotlin-testcontainers-and-spring-boot-native-537f624da895>
 - <https://github.com/testcontainers/testcontainers-java/blob/main/examples/sftp/src/test/java/org/example/SftpContainerTest.java>
- Exposed:
 - DAO: <https://ktor.io/docs/server-integrate-database.html#create-mapping>
- Cucumber:
 - <https://www.baeldung.com/java-cucumber-hooks>
 - <https://towardsdev.com/how-to-perfectly-test-your-ktor-application-94fb92c5a303>
 - <https://www.baeldung.com/cucumber-data-tables>
 - <https://github.com/cucumber/cucumber-jvm/tree/main/datatable>